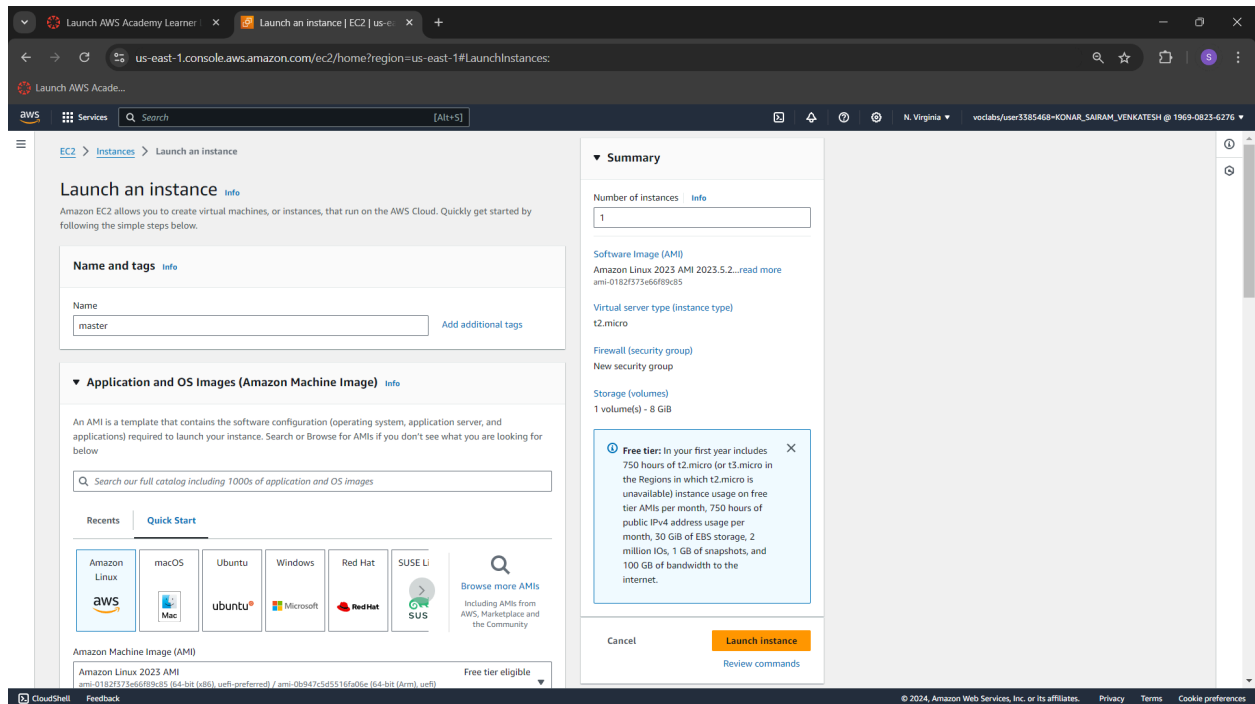


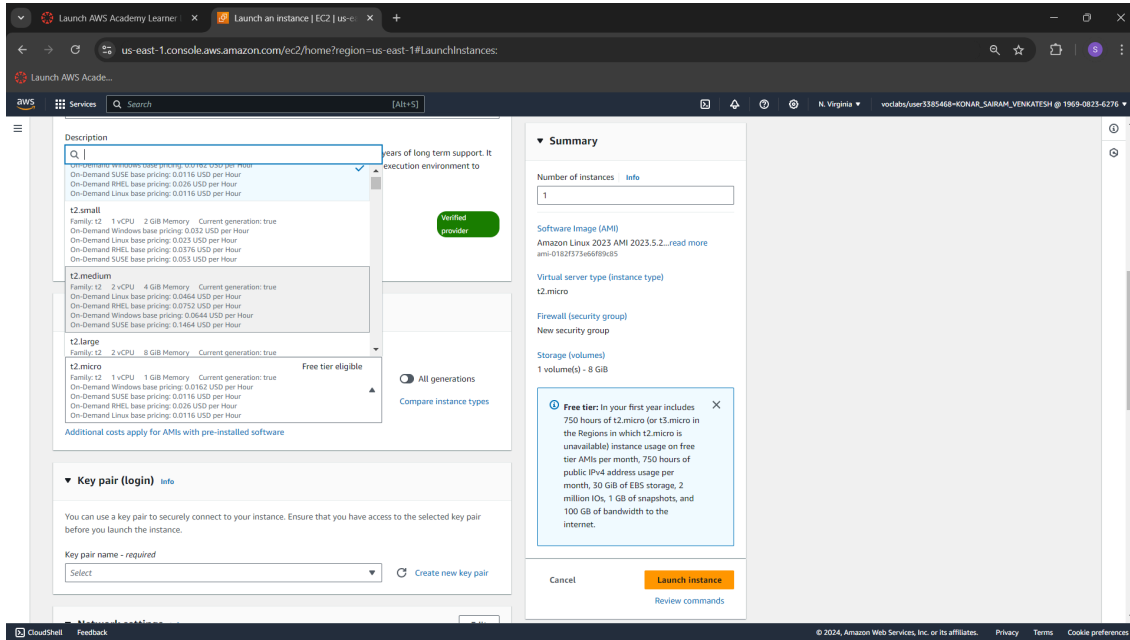
Aim: To understand the Kubernetes Cluster Architecture, install and Spin Up a Kubernetes Cluster on Linux Machines/Cloud Platforms.

Step 1: Set Up EC2 Instances.

- 1) Set up 3 EC2 instances called master, node1, node2
Select Amazon Linux as the OS image.



IMPORTANT: The default instance type and free one provided by AWS is t2.micro, which provides only 1CPU and 1 GiB of memory. For running Kubernetes, a minimum of 2 CPUs and 2GiB of RAM is required, hence change **t2.micro** to **t2.medium**.



2) Select a key pair, it may be the default (vockey provided by AWS Academy) or you may create one.

These are the 3 instances created. Click on the instance ID of each.

Instances (3) Info										
Find Instance by attribute or tag (case-sensitive)										
All states										
	Name	Instance ID	Instance state	Instance type	Status check	Alarm status	Availability Zone	Public IPv4 DNS	Public IPv4	Elastic IP
<input type="checkbox"/>	master	i-0619c0876a49226f6	Running	t2.medium	Initializing	View alarms	us-east-1c	ec2-54-91-165-28.com...	54.91.165.28	-
<input type="checkbox"/>	node2	i-078ec59bb0bf94b4e	Running	t2.medium	Initializing	View alarms	us-east-1c	ec2-54-165-218-135.co...	54.165.218.135	-
<input type="checkbox"/>	node1	i-03f7f63352d4f4e10	Running	t2.medium	Initializing	View alarms	us-east-1c	ec2-54-167-96-146.co...	54.167.96.146	-

3) Click on connect.

EC2 > Instances > i-0619c0876a49226f6										
Instance summary for i-0619c0876a49226f6 (master) Info										
Updated less than a minute ago										
Instance ID i-0619c0876a49226f6 (master)			Public IPv4 address 54.91.165.28 open address			Private IPv4 addresses 172.31.24.88				
IPv6 address -			Instance state Running			Public IPv4 DNS ec2-54-91-165-28.compute-1.amazonaws.com open address				
Hostname type IP name: ip-172-31-24-88.ec2.internal			Private IP DNS name (IPv4 only) ip-172-31-24-88.ec2.internal			Elastic IP addresses -				
Answer private resource DNS name IPv4 (A)			Instance type t2.medium			AWS Compute Optimizer finding Opt-in to AWS Compute Optimizer for recommendations. Learn more				
Auto-assigned IP address 54.91.165.28 [Public IP]			VPC ID vpc-05537c9b6c7862468			Auto Scaling Group name -				
IAM Role -			Subnet ID subnet-f74da678e35a8b10							

4) Maintain the default options and click on connect.

EC2 > Instances > i-0619c0876a49226f6 > Connect to instance

Connect to instance info


Connect to your instance i-0619c0876a49226f6 (master) using any of these options

EC2 Instance Connect

Session Manager

SSH client

EC2 serial console

**Port 22 (SSH) is open to all IPv4 addresses**
Port 22 (SSH) is currently open to all IPv4 addresses, indicated by 0.0.0.0/0 in the inbound rule in [your security group](#). For increased security, consider restricting access to only the EC2 Instance Connect service IP addresses for your Region: 18.206.107.24/29. [Learn more](#).

Instance ID
i-0619c0876a49226f6 (master)


Connection Type

☒ Connect using EC2 Instance Connect
Connect using the EC2 Instance Connect browser-based client, with a public IPv4 address.

☐ Connect using EC2 Instance Connect Endpoint
Connect using the EC2 Instance Connect browser-based client, with a private IPv4 address and a VPC endpoint.

Public IPv4 address
54.91.165.28




Username
Enter the username defined in the AMI used to launch the instance. If you didn't define a custom username, use the default username, ec2-user.

**Note:** In most cases, the default username, ec2-user, is correct. However, read your AMI usage instructions to check if the AMI owner has changed the default AMI username.

Cancel

Connect

5) After doing the above steps, the terminal for all 3 nodes can be seen.



PERFORM THE FOLLOWING STEPS ON ALL 3 MACHINES

Step 2: Installation of Docker

1) Use command

`'sudo su'`

This allows you to act as the root user of the terminal

```
_____/m/_____  
[ec2-user@ip-172-31-21-124 ~]$ sudo su  
[root@ip-172-31-21-124 ec2-user]#
```

2) We can install docker using YUM(Yellowdog Updater, Modified). Use the command

`'yum install docker -y'`

```
[root@ip-172-31-21-124 ec2-user]# yum install docker -y  
Last metadata expiration check: 0:05:01 ago on Fri Sep 13 16:47:44 2024.  
Dependencies resolved.  
=====
```

Package	Size	Architecture	Version	Repository
---------	------	--------------	---------	------------

```
-----  
Installing:  
docker                                44 M                x86_64                25.0.6-1.amzn2023.0.2    amazonlinux  
Installing dependencies:  
containerd                            35 M                x86_64                1.7.20-1.amzn2023.0.1    amazonlinux  
iptables-libs                         401 k               x86_64                1.8.8-3.amzn2023.0.2    amazonlinux  
iptables-nft                         193 k               x86_64                1.8.8-3.amzn2023.0.2    amazonlinux  
libcgroup                             75 k                x86_64                3.0-1.amzn2023.0.1      amazonlinux  
libnetfilter_conntrack               58 k                x86_64                1.0.8-2.amzn2023.0.2    amazonlinux  
libnftnl                             30 k                x86_64                1.0.1-19.amzn2023.0.2    amazonlinux  
libnftnl                             84 k                x86_64                1.2.2-2.amzn2023.0.2    amazonlinux  
pigz                                  83 k                x86_64                2.5-1.amzn2023.0.3      amazonlinux  
runc                                  83 k                x86_64                1.1.13-1.amzn2023.0.1    amazonlinux  
Total: 3.2 M  
Transaction Summary  
-----  
Install 10 Packages  
Total download size: 84 M  
Installed size: 317 M  
Downloading Packages:  
(1/10): iptables-libs-1.8.8-3.amzn2023.0.2.x86_64.rpm 4.0 MB  
/s | 401 kB 00:00  
(2/10): iptables-nft-1.8.8-3.amzn2023.0.2.x86_64.rpm 3.1 MB  
/s | 193 kB 00:00  
(3/10): libcgroup-3.0-1.amzn2023.0.1.x86_64.rpm 2.1 MB  
/s | 75 kB 00:00  
(4/10): libnetfilter_conntrack-1.0.8-2.amzn2023.0.2.x86_64.rpm 2.9 MB  
/s | 58 kB 00:00  
(5/10): libnftnl-1.0.1-19.amzn2023.0.2.x86_64.rpm 685 kB  
/s | 30 kB 00:00  
(6/10): libnftnl-1.2.2-2.amzn2023.0.2.x86_64.rpm 1.5 MB  
/s | 84 kB 00:00  
(7/10): pigz-2.5-1.amzn2023.0.3.x86_64.rpm 2.4 MB  
/s | 83 kB 00:00  
(8/10): runc-1.1.13-1.amzn2023.0.1.x86_64.rpm 23 MB  
/s | 3.2 MB 00:00  
(9/10): docker-25.0.6-1.amzn2023.0.2.x86_64.rpm 35 MB  
/s | 44 MB 00:01  
(10/10): containerd-1.7.20-1.amzn2023.0.1.x86_64.rpm 22 MB  
/s | 35 MB 00:01  
-----  
Total  
/s | 84 MB 00:01  
Running transaction check  
Transaction check succeeded.  
Running transaction test  
Transaction test succeeded.
```

```

Running transaction
Preparing      : runc-1.1.13-1.amzn2023.0.1.x86_64                                1/10
Installing     : containerd-1.7.20-1.amzn2023.0.1.x86_64                        1/10
Installing     : containerd-1.7.20-1.amzn2023.0.1.x86_64                        2/10
Running scriptlet: containerd-1.7.20-1.amzn2023.0.1.x86_64                      2/10
Installing     : pigz-2.5-1.amzn2023.0.3.x86_64                                3/10
Installing     : libnftnl-1.2.2-2.amzn2023.0.2.x86_64                         4/10
Installing     : libnftnl-1.2.2-2.amzn2023.0.2.x86_64                         5/10
Installing     : libnftfilter_contrack-1.0.8-2.amzn2023.0.2.x86_64              6/10
Installing     : iptables-libse-1.8.8-3.amzn2023.0.2.x86_64                   7/10
Installing     : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64                     8/10
Running scriptlet: iptables-nft-1.8.8-3.amzn2023.0.2.x86_64                   8/10
Installing     : libcgroupr-3.0-1.amzn2023.0.1.x86_64                         9/10
Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64                      10/10
Installing     : docker-25.0.6-1.amzn2023.0.2.x86_64                         10/10
Running scriptlet: docker-25.0.6-1.amzn2023.0.2.x86_64                       10/10
Created symlink /etc/systemd/system/sockets.target.wants/docker.socket → /usr/lib/systemd/system/docker.socket.

Verifying      : containerd-1.7.20-1.amzn2023.0.1.x86_64                                1/10
Verifying      : docker-25.0.6-1.amzn2023.0.2.x86_64                                2/10
Verifying      : iptables-libse-1.8.8-3.amzn2023.0.2.x86_64                     3/10
Verifying      : iptables-nft-1.8.8-3.amzn2023.0.2.x86_64                     4/10
Verifying      : libcgroupr-3.0-1.amzn2023.0.1.x86_64                         5/10
Verifying      : libnftfilter_contrack-1.0.8-2.amzn2023.0.2.x86_64             6/10
Verifying      : libnftnl-1.2.2-2.amzn2023.0.2.x86_64                         7/10
Verifying      : libnftnl-1.2.2-2.amzn2023.0.2.x86_64                         8/10
Verifying      : pigz-2.5-1.amzn2023.0.3.x86_64                                9/10
Verifying      : runc-1.1.13-1.amzn2023.0.1.x86_64                             10/10

Installed:
  containerd-1.7.20-1.amzn2023.0.1.x86_64      docker-25.0.6-1.amzn2023.0.2.x86_64      iptables-libse-1.8.8-3.amzn2023.0.2.x86_64      iptables-nft-1.8.8-3.amzn2023.0.2.x86_64
  libcgroupr-3.0-1.amzn2023.0.1.x86_64      libnftfilter_contrack-1.0.8-2.amzn2023.0.2.x86_64      libnftnl-1.0.1-19.amzn2023.0.2.x86_64
  pigz-2.5-1.amzn2023.0.3.x86_64          runc-1.1.13-1.amzn2023.0.1.x86_64      libnftnl-1.2.2-2.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-21-124 ec2-user]#

```

3) Now, configure a daemon.json file by using the following chain of commands.

- `cd /etc/docker`
- `cat <<EOF | sudo tee /etc/docker/daemon.json`

```

{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF

```
- `sudo systemctl enable docker`
- `sudo systemctl daemon-reload`
- `sudo systemctl restart docker`

```

[root@ip-172-31-20-253 ec2-user]# cd /etc/docker
cat <<EOF | sudo tee /etc/docker/daemon.json
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
EOF
sudo systemctl enable docker
sudo systemctl daemon-reload
sudo systemctl restart docker
{
  "exec-opts": ["native.cgroupdriver=systemd"],
  "log-driver": "json-file",
  "log-opts": {
    "max-size": "100m"
  },
  "storage-driver": "overlay2"
}
Created symlink /etc/systemd/system/multi-user.target.wants/docker.service → /usr/lib/systemd/system/docker.service.
[root@ip-172-31-20-253 ec2-user]#

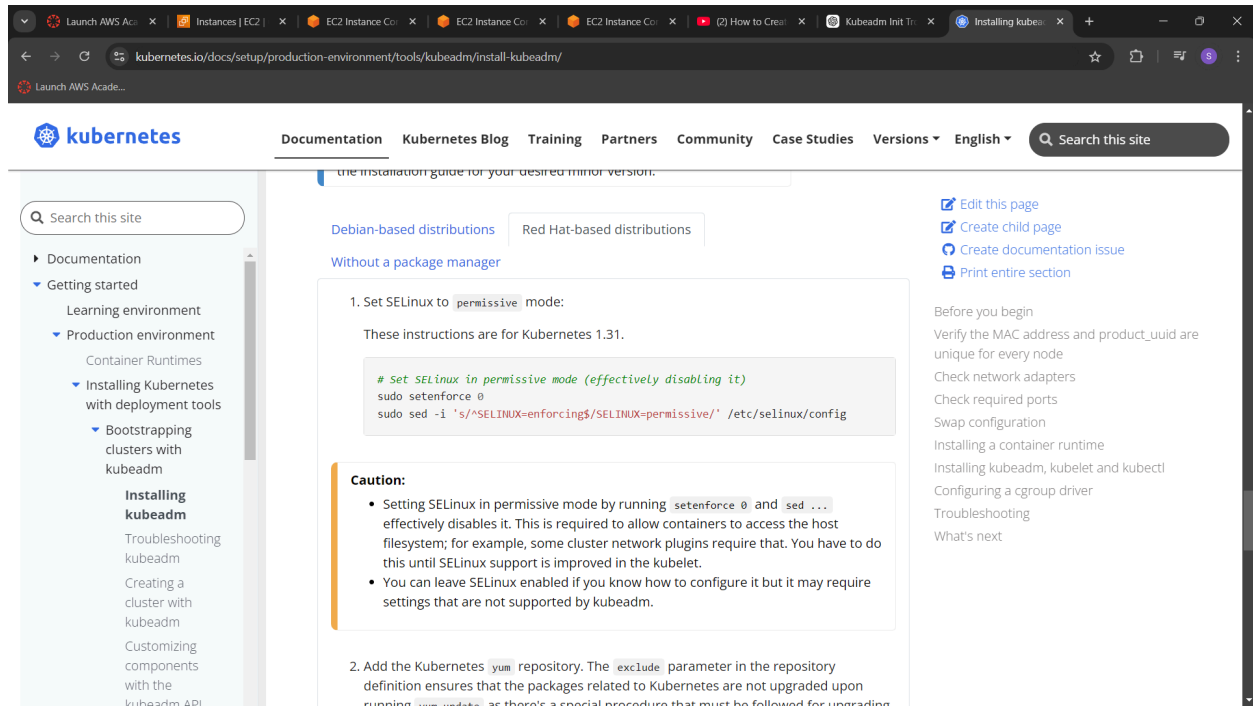
```

Step 3: Installing Kubernetes

- 1) For installing kubernetes, we will be using kubeadm, a framework used for creating kubernetes clusters using command line.

<https://kubernetes.io/docs/setup/production-environment/tools/kubeadm/install-kubeadm/>

The following will be visible when you visit the website.



- 2) Select **red hat-based distributions** as amazon linux is based on red hat.

```
sudo setenforce 0
```

→ sets SELinux to permissive mode

```
sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config
```

→ edits the SELinux configuration file (/etc/selinux/config) to make the change persistent across reboots. If not used, SELinux reverts to enforcing mode after reboot.

Setting SELinux to permissive mode during Kubernetes installation prevents permission-related issues with container runtimes and components that may not function correctly under SELinux's enforcing policies.

Run the following commands:

- `sudo setenforce 0`
- `sudo sed -i 's/^SELINUX=enforcing$/SELINUX=permissive/' /etc/selinux/config`

- `cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo`
`[kubernetes]`
`name=Kubernetes`
`baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/`
`enabled=1`
`gpgcheck=1`
`gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key`
`exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni`
`EOF`

This comamnd is a repository script to create a kubernetes repository

```
[root@ip-172-31-21-124 ec2-user]# cat <<EOF | sudo tee /etc/yum.repos.d/kubernetes.repo
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
EOF
[kubernetes]
name=Kubernetes
baseurl=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/
enabled=1
gpgcheck=1
gpgkey=https://pkgs.k8s.io/core:/stable:/v1.31/rpm/repodata/repomd.xml.key
exclude=kubelet kubeadm kubectl cri-tools kubernetes-cni
```

- yum repolist

This command shows the repositories created on the machine.

```
[root@ip-172-31-21-124 ec2-user]# yum repolist
```

repo id	repo name
amazonlinux	Amazon Linux 2023 repository
kernel-livepatch	Amazon Linux 2023 Kernel Livepatch repository
kubernetes	Kubernetes

```
[root@ip-172-31-21-124 ec2-user]#
```

Next step is to install kubelet, kubeadm, kubectl

- `sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes`

```
[root@ip-172-31-21-124 ec2-user]# sudo yum install -y kubelet kubeadm kubectl --disableexcludes=kubernetes
kubernetes
Kubernetes
56 kB/s | 9.4 kB
00
Dependencies resolved.
=====
Package                               Architecture      Version           Repository
-----
Installing:
kubeadm                               x86_64            1.31.1-150500.1.1  kubernetes
11 M
kubectl                               x86_64            1.31.1-150500.1.1  kubernetes
11 M
kubelet                               x86_64            1.31.1-150500.1.1  kubernetes
15 M
Installing dependencies:
conntrack-tools                       x86_64            1.4.6-2.amzn2023.0.2  amazonlinux
208 k
cri-tools                             x86_64            1.31.1-150500.1.1  kubernetes
6.9 M
kubernetes-cni                       x86_64            1.5.1-150500.1.1  kubernetes
7.1 M
libnetfilter_cthelper                x86_64            1.0.0-21.amzn2023.0.2  amazonlinux
24 k
libnetfilter_cttimeout               x86_64            1.0.0-19.amzn2023.0.2  amazonlinux
24 k
libnetfilter_queue                   x86_64            1.0.5-2.amzn2023.0.2  amazonlinux
30 k
Transaction Summary
=====
```

```

Running scriptlet: kubect1-1.31.1-150500.1.1.x86_64
9/9
Verifying      : contrack-tools-1.4.6-2.amzn2023.0.2.x86_64
1/9
Verifying      : libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64
2/9
Verifying      : libnetfilter_cttimeout-1.0.0-19.amzn2023.0.2.x86_64
3/9
Verifying      : libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64
4/9
Verifying      : cri-tools-1.31.1-150500.1.1.x86_64
5/9
Verifying      : kubeadm-1.31.1-150500.1.1.x86_64
6/9
Verifying      : kubect1-1.31.1-150500.1.1.x86_64
7/9
Verifying      : kubelet-1.31.1-150500.1.1.x86_64
8/9
Verifying      : kubernetes-cni-1.5.1-150500.1.1.x86_64
9/9

Installed:
contrack-tools-1.4.6-2.amzn2023.0.2.x86_64      cri-tools-1.31.1-150500.1.1.x86_64      kubeadm-1.31.1-150500.1.1.x86_64      kubect1-1.31.1-150500.1.1.x86_
64
kubelet-1.31.1-150500.1.1.x86_64      kubernetes-cni-1.5.1-150500.1.1.x86_64      libnetfilter_cthelper-1.0.0-21.amzn2023.0.2.x86_64      libnetfilter_cttimeout-1.0.0-1
9.amzn2023.0.2.x86_64
libnetfilter_queue-1.0.5-2.amzn2023.0.2.x86_64

Complete!
[root@ip-172-31-21-124 ec2-user]#

```

Now, we need to enable the kubelet service. Run the command

- `sudo systemctl enable --now kubelet`

```

[root@ip-172-31-21-124 ec2-user]# sudo systemctl enable --now kubelet
Created symlink /etc/systemd/system/multi-user.target.wants/kubelet.service → /usr/lib/systemd/system/kubelet.service.

```

PERFORM THE FOLLOWING ON ONLY THE MASTER MACHINE

- 1) Firstly, we need to initialize kubernetes. For this, run the command:

- `kubeadm init`

```

[root@ip-172-31-21-124 ec2-user]# kubeadm init
[init] Using Kubernetes version: v1.31.0
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
[WARNING FileExisting-tc]: tc not found in system path
[preflight] Pulling images required for setting up a Kubernetes cluster
[preflight] This might take a minute or two, depending on the speed of your internet connection
[preflight] You can also perform this action beforehand using 'kubeadm config images pull'
W0119:05:14:24:28.45614 checks.go:846] detected that the sandbox image "registry.k8s.io/pause:3.8" of the container runtime is inconsistent with that used by kubeadm. It is recommended to use "registry.k8s.io/pause:3.10" as the CNI sandbox image.
[certs] Using certificateDir folder "/etc/kubernetes/pki"
[certs] Generating "ca" certificate and key
[certs] Generating "apiserver" certificate and key
[certs] apiserver serving cert is signed for DNS names [ip-172-31-21-124.ec2.internal kubernetes.kubernetes.default kubernetes.default.svc kubernetes.default.svc.cluster.local] and IPs [10.96.0.1 172.31.21.124]
[certs] Generating "apiserver-kubelet-client" certificate and key
[certs] Generating "front-proxy-ca" certificate and key
[certs] Generating "front-proxy-client" certificate and key
[certs] Generating "etcd/ca" certificate and key
[certs] Generating "etcd/server" certificate and key
[certs] etcd/server serving cert is signed for DNS names [ip-172-31-21-124.ec2.internal localhost] and IPs [172.31.21.124 127.0.0.1 ::1]
[certs] Generating "etcd/peer" certificate and key
[certs] etcd/peer serving cert is signed for DNS names [ip-172-31-21-124.ec2.internal localhost] and IPs [172.31.21.124 127.0.0.1 ::1]
[certs] Generating "etcd/healthcheck-client" certificate and key
[certs] Generating "apiserver-etcd-client" certificate and key
[certs] Generating "sa" key and public key
[kubeconfig] Using kubeconfig folder "/etc/kubernetes"
[kubeconfig] Writing "admin.conf" kubeconfig file
[kubeconfig] Writing "super-admin.conf" kubeconfig file
[kubeconfig] Writing "kubelet.conf" kubeconfig file
[kubeconfig] Writing "controller-manager.conf" kubeconfig file
[kubeconfig] Writing "scheduler.conf" kubeconfig file
[etcd] Creating static Pod manifest for local etcd in "/etc/kubernetes/manifests"
[control-plane] Using manifest folder "/etc/kubernetes/manifests"

```

```

[api-check] Waiting for a healthy API server. This can take up to 4m0s
[api-check] The API server is healthy after 5.002506688s
[upload-config] Storing the configuration used in ConfigMap "kubeadm-config" in the "kube-system" Namespace
[kubelet] Creating a ConfigMap "kubelet-config" in namespace kube-system with the configuration for the kubelets in the cluster
[upload-certs] Skipping phase. Please see --upload-certs
[mark-control-plane] Marking the node ip-172-31-21-124.ec2.internal as control-plane by adding the labels: [node-role.kubernetes.io/control-plane node.kubernetes.io/exclude-from-external-load-balancers]
[mark-control-plane] Marking the node ip-172-31-21-124.ec2.internal as control-plane by adding the taints [node-role.kubernetes.io/control-plane:NoSchedule]
[bootstrap-token] Using token: 76rlql.gpqk9pmv8t9osgg4
[bootstrap-token] Configuring bootstrap tokens, cluster-info ConfigMap, RBAC Roles
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to get nodes
[bootstrap-token] Configured RBAC rules to allow Node Bootstrap tokens to post CSRs in order for nodes to get long term certificate credentials
[bootstrap-token] Configured RBAC rules to allow the csrapprover controller automatically approve CSRs from a Node Bootstrap Token
[bootstrap-token] Configured RBAC rules to allow certificate rotation for all node client certificates in the cluster
[bootstrap-token] Creating the "cluster-info" ConfigMap in the "kube-public" namespace
[kubelet-finalize] Updating "/etc/kubernetes/kubelet.conf" to point to a rotatable kubelet client certificate and key
[addons] Applied essential addon: CoreDNS
[addons] Applied essential addon: kube-proxy

Your Kubernetes control-plane has initialized successfully!

To start using your cluster, you need to run the following as a regular user:

  mkdir -p $HOME/.kube
  sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
  sudo chown $(id -u):$(id -g) $HOME/.kube/config

Alternatively, if you are the root user, you can run:

  export KUBECONFIG=/etc/kubernetes/admin.conf

You should now deploy a pod network to the cluster.
Run "kubect1 apply -f [podnetwork].yaml" with one of the options listed at:
  https://kubernetes.io/docs/concepts/cluster-administration/addons/

Then you can join any number of worker nodes by running the following on each as root:

kubeadm join 172.31.21.124:6443 --token 76rlql.gpqk9pmv8t9osgg4 \
  --discovery-token-ca-cert-hash sha256:a031677cdc93a202c5f89e2370eb7231d212211bac6f816cddcb547b7f00a4f2
[root@ip-172-31-21-124 ec2-user]#

```


From the output, we will receive a command that is used to link the nodes to the master. Copy it and save it somewhere local.

```
kubeadm join 172.31.21.124:6443 --token 76r1ql.qpgk9pmv8t9osgg4 \
--discovery-token-ca-cert-hash sha256:a03f677cdc93d202c5f89e2370eb7231d212211bac6f816cddcb547b7f00a4f2
```

2) From the output, we receive the following commands:

- `mkdir -p $HOME/.kube`
- `sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config`
- `sudo chown $(id -u):$(id -g) $HOME/.kube/config`

Run these commands.

To start using your cluster, you need to run the following as a regular user:

```
mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

```
[root@ip-172-31-21-124 ec2-user]# mkdir -p $HOME/.kube
sudo cp -i /etc/kubernetes/admin.conf $HOME/.kube/config
sudo chown $(id -u):$(id -g) $HOME/.kube/config
```

3) To check whether nodes are connected, run the command

- `kubectl get nodes`

This output shows only master is connected right now.

```
[root@ip-172-31-20-253 docker]# kubectl get nodes
NAME                                STATUS    ROLES    AGE    VERSION
ip-172-31-20-253.ec2.internal      NotReady  control-plane  2m19s  v1.31.1
```

PERFORM THE FOLLOWING ONLY ON THE NODE MACHINES

Use the command that you had copied before and use them on the node machines.

```
[root@ip-172-31-23-39 docker]# kubeadm join 172.31.21.12:6443 --token 5mqobq.tz0ofaxdr0jhd384 --discovery-token-ca-cert-hash sha256:c7f71cc84c1dd813b0ea20de72a188dab31a490d2aceealc30d9744f5e1e7f59
[preflight] Running pre-flight checks
[WARNING FileExisting-socat]: socat not found in system path
error execution phase preflight: couldn't validate the identity of the API Server: failed to request the cluster-info ConfigMap: Get "https://172.31.21.12:6443/api/v1/namespaces/kube-public/configmaps/cluster-info?timeout=10s": context deadline exceeded
To see the stack trace of this error execute with --v=5 or higher
[root@ip-172-31-23-39 docker]#
```

ERROR

Here, an error occurs as 'Context Deadline Exceeded'. Hence, the final output is not obtained.

Conclusion:

In this experiment, we have learned how to create kubernetes clusters on a linux terminal, how the ssh command works on a local terminal and what requirements are necessary to create kubernetes clusters. In the final step, while working on the EC2 instance terminal, the connection of the master node to the slave nodes does not happen as it gives an error of 'Context Deadline Exceeded'.